

# Multi-Agent Reinforcement Learning for Coordinated Block Transportation

Maikel Withagen (S1867733)<sup>a</sup>      Steven Bosch (S1861948)<sup>a</sup>  
Thomas Derksen (S2016346)<sup>a</sup>

<sup>a</sup> *University of Groningen, Faculty of Mathematics and Natural Sciences,  
Nijenborgh 9, 9747 AG, Groningen*

## Abstract

## 1 Introduction

One of the returning issues in machine learning is that it is often hard to decide which learning algorithm should be used for a certain application. Over the years, a lot of valid learning algorithms have emerged, each with their own advantages and disadvantages. Some algorithms might be faster, while others might be more precise, and an algorithm that is better in one area usually performs worse in another. Matching the right task with the right algorithm is not always a straightforward task.

In this paper we will look at the problem of multi-agent box-transportation. For this task, two virtual robots will have to cooperate to transport a box to a certain goal area, while avoiding certain obstacles. This is a problem that has already garnered some attention in machine learning research. In 2006, Wang and de Silva tried to solve this problem with reinforcement learning techniques [4]. They compared the performance of a Q-learning algorithm with the performance of team Q-learning in a dynamic environment. In 2010, a team at the University of Deft also showed how you could use Q-learning, team Q-learning and the WoLF-PHC algorithm to accomplish a similar task in a static environment [1].

We will also compare the Q-learning and team Q-learning algorithm, since they seem most suitable for a task like this. For our initial task, two agents have to transport a box in a two dimensional static environment. We will compare the performance of our two reinforcement learning algorithms by looking at the time it takes for them to converge to the optimal solution (given that they find it).

Secondly, we will also expand the task by adding another goal area or adding additional agents. **However, we still have to see if this is doable.**


## 2 Application

The box-transportation problem is significant in both the theoretical and practical domain. From a theoretical point of view, it is a relatively simple problem with clear boundaries. Although there are multiple versions of this problem, the basic idea is intuitively easy to understand. Because of its simplicity, it is not hard to extend it to more complicated, real-life situations. It is also a good task to observe the effects of different learning algorithms since, in

the more simple, stripped down versions of this task, there are not a lot of parameters that influence the outcome. Finally, it is an interesting way of studying cooperation between multiple autonomous agents.

Although we use a simulation in this paper, the task has successfully been completed by physical robots. For example, in [2] a team of two Pioneer 2 mobile robots was used to physically push a box to a goal area while a third robot served as an observer. In [3] they took a different approach, where they used a swarm of smaller custom robots to collectively move a box to a designated spot (as can be seen in figure 1).

Figure 1: Robots from [3] locating a lit box



It is not hard to come up with practical applications for these kind of systems. In the future, they might for example be useful in construction, or to clear the rubble after a building has collapsed.

**Our extentions are also interesting, for obvious reasons.**

### 3 Method

To compare the various learning algorithms, we have written a multi-agent simulation in Python in which agents have to move an object to a certain goal area. In this section, the various components of the simulation will be explained.

#### 3.1 The Environment

The environment consists of an  $x$  by  $y$  grid in which the agents can move around. There are a few different types of cells:

**Free cells** These are cells that agents can freely move to. Once an agent or a block moves to a free cell, this cell becomes occupied. They are denoted by dots in the visual representation

of the simulation.

**Walls** These cells are occupied, so an agent can never move to them. They are initialized at the start of the simulation and will not change during the simulation. There are outer walls, which are on the edges of the grid, and some walls inside the grid to make the environment more complex. They look like hash tags in the visualization.

**Block** The block is an item that has to be transported to a certain goal area by the agents. It occupies one cell, and agents can't move through it. It can only be moved if it's grasped by two agents, and those agents move in the same direction. If the block reaches the goal, a run ends. In the simulation, it is denoted by a B.

**Goal** The goal cell is basically the same as a free cell, with the exception that if the block reaches the goal, a run of the simulation ends and a new run starts. Agents can freely move on and over the goal, just like with free cells. The goal is represented by a B in the simulation.

**Agent** We'll talk about these in more detail in the next subsection. They occupy one cell, so other agents can't move through them. In the simulation they look like A's. In **FIGURE X** you can see how the visual representation of the simulation looks.

### 3.2 The Agents

For our initial simulation, we start out with two agents. At the start of a run, the agents are spawned at their respective start locations which are defined in the algorithm. Each step, an agent can perform one of five actions:

- *Stay*
- *Left*
- *Up*
- *Right*
- *Down*
- *Grab*

The actions mostly speak for themselves. The agent can stay where it is, move in one of four directions or grab the block. The *grab* action can be performed at any step, but will only do something when an agent is next to the block. Once an agent has grabbed the block it will not let go. The move actions *left*, *up*, *right* and *down* only do something when the cell the agent wants to move to is actually free.

Each step in the simulation, an agent chooses the best action based on its learning algorithm. Then, the agents perform their actions and update the values related to their learning algorithms (usually these are the values in their Q-tables). This process repeats itself until the block has reached the goal. After that, a new run starts and the agents are reset to their starting positions.

### 3.3 The Learning Algorithms

For this paper, we compared **TWO machine learning algorithms**: Q-learning and team Q-learning.

## 4 Results

## 5 Discussion

## References

- [1] Lucian Busoniu, Robert Babuska, and Bart de Schutter. Multi-agent reinforcement learning: An overview. In Dipti Srinivasan and Lakhmi Jain, editors, *Innovations in Multi-Agent Systems and Applications*, chapter 1, pages 183–221. Springer Berlin Heidelberg, Berlin, 2010.
- [2] Brian P. Gerkey and Maja J. Mataric. Pusher-watcher: an approach to fault-tolerant tightly-coupled robot coordination. In *International Conference on Robotics and Automation*, pages 464–469, 2002.
- [3] C. Ronald Kube and Hon Zhang. The use of perceptual cues in multi-robot box-pushing. In *International Conference on Robotics and Automation*, pages 2085 – 2090, 1996.
- [4] Ying Wang and Clarence de Silva. Multi-robot box-pushing: Single-agent q-learning vs. team q-learning. In *International Conference on Intelligent Robots and Systems*, 2006.